

solrpres

March 2, 2015

1 A quick tour of Lucene and Solr

[Lucene](#) is a widely used information retrieval library in Java.

- implements the TF-IDF model
- simple model for indexing and search
- easy to learn, fast, and scalable
- written by Doug Cutting, who later co-wrote Hadoop
- I've been using it since 2001 ([DSpace](#), [Canary](#), [WDL](#), [Chronam](#), [GW Libs](#))
- free, open source software

1.1 Lucene basics

- everything is a Document (source format independent)
- Documents have Fields
- Fields can be tokenized, indexed, stored, multivalued (or not)
- text is Tokenized, Analyzed on index and query
- queries are parsed, then Tokenized, Analyzed, and results ranked

See [a simple example of code using the Lucene API](#) and also [more extensive Lucene documentation](#).

1.2 Solr basics

[Solr](#) is a search application (or server) that makes Lucene even easier to implement and scale.

- Lucene under the hood
- simple “index anything” or schema-based configuration
- simple web API to index and search
- scales to many servers
- very widely used
- free, open source software

2 Working with Solr

- index and search using web API - easier code
- define a schema first, or don't
- csv, json, xml results easy to parse in any language
- handles many technical details (caching, faceting, spelling corrections, clustering, hit highlighting)
- the bee's knees, basically

2.1 Searching solr

It's as easy as sending an HTTP request. Let's use `python-requests` to handle HTTP.

```
In [1]: import requests
        req = requests.get("http://localhost:8983/solr/gettingstarted/select?wt=json&q=foundation")
        resp = req.json()
        resp.keys()
```

```
Out[1]: [u'responseHeader', u'response']
```

The response header is pretty straightforward:

```
In [2]: resp['responseHeader']
```

```
Out[2]: {u'QTime': 38,
        u'params': {u'indent': u'true', u'q': u'foundation', u'wt': u'json'},
        u'status': 0}
```

```
In [3]: resp['response'].keys()
```

```
Out[3]: [u'start', u'maxScore', u'numFound', u'docs']
```

```
In [6]: resp['response']['start']
```

```
Out[6]: 0
```

```
In [7]: resp['response']['maxScore']
```

```
Out[7]: 0.37489003
```

```
In [8]: resp['response']['numFound']
```

```
Out[8]: 3106
```

That's the basic header. We're starting with the first page of results, and it's a zero-based index, which explains why "start" is 0. We'll look more at scores in a moment.

Let's take a look at the results first.

```
In [10]: len(resp['response']['docs'])
```

```
Out[10]: 10
```

```
In [11]: docs = resp['response']['docs']
        docs[0]
```

```
Out[11]: {u'_version_': 1494510828527812608,
        u'author': [u'Isaac Asimov'],
        u'cat': [u'book'],
        u'genre_s': u'scifi',
        u'id': u'0553293354',
        u'inStock': [True],
        u'name': [u'Foundation'],
        u'price': [7.99],
        u'sequence_i': 1,
        u'series_t': [u'Foundation Novels']}
```

```
In [12]: docs[1]
```

```

Out[12]: {u'_version_': 1494510737274437632,
          u'cat': [u'software', u'search'],
          u'features': [u'No accents here',
                        u'This is an e acute: \xe9',
                        u'eaoui with circumflexes: \xea\xe2\xee\xfa\xfb',
                        u'eaoui with umlauts: \xeb\xe4\xef\xfc',
                        u'tag with escaped chars: <nicetag/>',
                        u'escaped ampersand: Bonnie & Clyde',
                        u'Outside the BMP:\U00010308 codepoint=10308, a circle with an x inside. UTF8=f0908c88 UTF16='],
          u'id': u'UTF8TEST',
          u'inStock': [True],
          u'manu': [u'Apache Software Foundation'],
          u'name': [u'Test with some UTF-8 encoded characters'],
          u'price': [0.0]}

```

```

In [21]: for doc in docs:
          print "%s: %s, $%s" % (doc['id'], doc.get('name', '(no title)'), doc.get('price', '-'))

```

```

0553293354: [u'Foundation'], $[7.99]
UTF8TEST: [u'Test with some UTF-8 encoded characters'], $[0.0]
SOLR1000: [u'Solr, the Enterprise Search Server'], $[0.0]
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-analytics/deprecated-list.html: (no title), $-
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clustering/carrot2/
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/logging/log4j/package-use.html: (
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/spelling/suggest/tst/package-use.
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-dataimporthandler-extras/deprecated-list.html: (no tit
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-solrj/org/apache/solr/client/solrj/util/package-use.ht
/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-test-framework/org/apache/solr/analysis/package-use.ht

```

Let's revisit this query and add relevance rank scores to the results.

```

In [22]: req = requests.get("http://localhost:8983/solr/gettingstarted/select?wt=json&indent=true&q=fou
          resp = req.json()
          docs = resp['response']['docs']
          docs[0]

```

```

Out[22]: {u'_version_': 1494510828527812608,
          u'author': [u'Isaac Asimov'],
          u'cat': [u'book'],
          u'genre.s': u'scifi',
          u'id': u'0553293354',
          u'inStock': [True],
          u'name': [u'Foundation'],
          u'price': [7.99],
          u'score': 0.37489003,
          u'sequence_i': 1,
          u'series_t': [u'Foundation Novels']}

```

```

In [24]: for doc in docs:
          print "%s - %s" % (doc['score'], doc.get('name', '(no title)'))

```

```

0.37489003 - [u'Foundation']
0.13254364 - [u'Test with some UTF-8 encoded characters']
0.115975685 - [u'Solr, the Enterprise Search Server']
0.05857657 - (no title)
0.05857657 - (no title)

```

```
0.05857657 - (no title)
0.05857657 - (no title)
0.05857657 - (no title)
0.05857657 - (no title)
0.05857657 - (no title)
```

3 Under the hood

Now let's take a closer look at how those scores are generated. Solr offers easy tools for examining a live index and debugging information about results.

```
In [25]: # Add parameter debugQuery=true, repeat
        req = requests.get("http://localhost:8983/solr/gettingstarted/select?wt=json&indent=true&q=foundation")
        resp = req.json()
        resp.keys()
```

```
Out[25]: [u'debug', u'reponseHeader', u'response']
```

```
In [26]: debug = resp['debug']
        debug.keys()
```

```
Out[26]: [u'parsedquery',
          u'track',
          u'explain',
          u'querystring',
          u'rawquerystring',
          u'parsedquery_toString',
          u'QParser',
          u'timing']
```

First, we can see that the query is parsed to determine how it should be processed. In the absence of named fields, it assumes a default, here `_text`.

```
In [30]: debug['QParser']
```

```
Out[30]: u'LuceneQParser'
```

```
In [27]: debug['rawquerystring']
```

```
Out[27]: u'foundation'
```

```
In [29]: debug['parsedquery']
```

```
Out[29]: u'_text: foundation'
```

The `timing` element gives the processing time in milliseconds. Note that a lot of tasks were not performed.

```
In [31]: debug['timing']
```

```
Out[31]: {u'prepare': {u'debug': {u'time': 0.0},
          u'expand': {u'time': 0.0},
          u'facet': {u'time': 0.0},
          u'highlight': {u'time': 0.0},
          u'mlt': {u'time': 0.0},
          u'query': {u'time': 4.0},
          u'stats': {u'time': 0.0},
          u'time': 4.0},
```

```

u'process': {u'debug': {u'time': 36.0},
u'expand': {u'time': 0.0},
u'facet': {u'time': 0.0},
u'highlight': {u'time': 0.0},
u'mlt': {u'time': 0.0},
u'query': {u'time': 6.0},
u'stats': {u'time': 0.0},
u'time': 42.0},
u'time': 46.0}

```

3.1 The good stuff: relevance scoring

Now we get to the good bits. The `explain` element contains the precise scoring details for each aspect of a parsed query.

```

In [34]: explain = debug['explain']
         explain.keys()

```

```

Out[34]: [u'0553293354',
u'UTF8TEST',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/spelling/suggest/tst/',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-solrj/org/apache/solr/client/solrj/util/pa',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/logging/log4j/package',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-dataimporthandler-extras/deprecated-list.h',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-test-framework/org/apache/solr/analysis/pa',
u'SOLR1000',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/cluster',
u'/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-analytics/deprecated-list.html']

```

```

In [36]: print explain['UTF8TEST']

```

```

0.13254364 = (MATCH) weight(_text:foundation in 16) [DefaultSimilarity], result of:
  0.13254364 = fieldWeight in 16, product of:
    1.0 = tf(freq=1.0), with freq of:
      1.0 = termFreq=1.0
    1.0603491 = idf(docFreq=1494, maxDocs=1588)
    0.125 = fieldNorm(doc=16)

```

```

In [37]: print explain['0553293354']

```

```

0.37489003 = (MATCH) weight(_text:foundation in 3) [DefaultSimilarity], result of:
  0.37489003 = fieldWeight in 3, product of:
    1.4142135 = tf(freq=2.0), with freq of:
      2.0 = termFreq=2.0
    1.0603491 = idf(docFreq=1494, maxDocs=1588)
    0.25 = fieldNorm(doc=3)

```

```

In [45]: req = requests.get("http://localhost:8983/solr/gettingstarted/select?wt=json&indent=true&q=chi")
         resp = req.json()
         resp['debug']['parsedquery']

```

```

Out[45]: u'_text:chinese name:tokenizer'

```

```

In [46]: for doc in resp['response']['docs']:
         print "%s - %s: %s" % (doc['score'], doc['id'], doc.get('name', '(no name)'))

```

```

0.10238882 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clustering/carrot2/
0.09220693 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clustering/carrot2/
0.009314307 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/changes/Changes.html: (no name)

```

Because this is a slightly more complex query, we can see a lot more going on in the explanations.

```

In [50]: explain = resp['debug']['explain']
         for id, result in explain.items():
             print "%s%s" % (id, result)

```

```

/Users/dchud/Downloads/apps/solr-5.0.0/docs/changes/Changes.html

```

```

0.009314307 = (MATCH) product of:
  0.018628614 = (MATCH) sum of:
    0.018628614 = (MATCH) weight(_text:chinese in 2) [DefaultSimilarity], result of:
      0.018628614 = score(doc=2,freq=4.0), product of:
        0.6558272 = queryWeight, product of:
          7.2716184 = idf(docFreq=2, maxDocs=1588)
          0.09018999 = queryNorm
        0.02840476 = fieldWeight in 2, product of:
          2.0 = tf(freq=4.0), with freq of:
            4.0 = termFreq=4.0
          7.2716184 = idf(docFreq=2, maxDocs=1588)
          0.001953125 = fieldNorm(doc=2)
    0.5 = coord(1/2)

```

```

/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clustering/carrot2/

```

```

0.09220693 = (MATCH) product of:
  0.18441387 = (MATCH) sum of:
    0.18441387 = (MATCH) weight(_text:chinese in 131) [DefaultSimilarity], result of:
      0.18441387 = score(doc=131,freq=2.0), product of:
        0.6558272 = queryWeight, product of:
          7.2716184 = idf(docFreq=2, maxDocs=1588)
          0.09018999 = queryNorm
        0.28119275 = fieldWeight in 131, product of:
          1.4142135 = tf(freq=2.0), with freq of:
            2.0 = termFreq=2.0
          7.2716184 = idf(docFreq=2, maxDocs=1588)
          0.02734375 = fieldNorm(doc=131)
    0.5 = coord(1/2)

```

```

/Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clustering/carrot2/

```

```

0.10238882 = (MATCH) product of:
  0.20477764 = (MATCH) sum of:
    0.20477764 = (MATCH) weight(_text:chinese in 122) [DefaultSimilarity], result of:
      0.20477764 = score(doc=122,freq=4.0), product of:
        0.6762287 = queryWeight, product of:
          7.75227 = idf(docFreq=1, maxDocs=1712)
          0.087229766 = queryNorm
        0.30282307 = fieldWeight in 122, product of:
          2.0 = tf(freq=4.0), with freq of:
            4.0 = termFreq=4.0
          7.75227 = idf(docFreq=1, maxDocs=1712)
          0.01953125 = fieldNorm(doc=122)
    0.5 = coord(1/2)

```

3.2 A better example

These examples are rather inscrutable; let's look at something more obvious. In addition to the example “getting started” API documentation indexed during the [Solr Tutorial](#) I've downloaded and indexed several hundred books from [Project Gutenberg](#). Let's try the “foundation” search again:

```
In [60]: req = requests.get("http://localhost:8983/solr/gettingstarted/select?q=foundation&wt=json&indent=2")
        resp = req.json()
        for doc in resp['response']['docs']:
            print "%s - %s: %s" % (doc['score'], doc['id'], doc.get('name', '(no name)'))

0.39567953 - 0553293354: [u'Foundation']
0.13989384 - UTF8TEST: [u'Test with some UTF-8 encoded characters']
0.122407116 - SOLR1000: [u'Solr, the Enterprise Search Server']
0.06483684 - /Volumes/PGCD0803/etext03/vbgle11h/index.html: (no name)
0.061824925 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/spelling/suggest/tst
0.061824925 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-dataimporthandler-extras/deprecated-list
0.061824925 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-solrj/org/apache/solr/client/solrj/util/Util
0.061824925 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-analytics/deprecated-list.html: (no name)
0.061824925 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clustering
0.061824925 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/logging/log4j/package-info.html
```

That first hit is new! Let's take a look.

```
In [62]: explain = resp['debug']['explain']
        print explain['0553293354']

0.39567953 = (MATCH) weight(_text:foundation in 1586) [DefaultSimilarity], result of:
  0.39567953 = fieldWeight in 1586, product of:
    1.4142135 = tf(freq=2.0), with freq of:
      2.0 = termFreq=2.0
    1.1191508 = idf(docFreq=1785, maxDocs=2012)
    0.25 = fieldNorm(doc=1586)

In [63]: print explain['SOLR1000']

0.122407116 = (MATCH) weight(_text:foundation in 1576) [DefaultSimilarity], result of:
  0.122407116 = fieldWeight in 1576, product of:
    1.0 = tf(freq=1.0), with freq of:
      1.0 = termFreq=1.0
    1.1191508 = idf(docFreq=1785, maxDocs=2012)
    0.109375 = fieldNorm(doc=1576)
```

Can you guess what it is?

```
In [64]: resp['response']['docs'][0]

Out[64]: {u'_version_': 1494510828527812608,
          u'author': [u'Isaac Asimov'],
          u'cat': [u'book'],
          u'genre_s': u'scifi',
          u'id': u'0553293354',
          u'inStock': [True],
          u'name': [u'Foundation'],
          u'price': [7.99],
          u'score': 0.39567953,
          u'sequence_i': 1,
          u'series_t': [u'Foundation Novels']}
```

This gives us a great opportunity to see how relevance calculations vary with specific queries. First, we add the bare term “asimov”, so our search query is now “foundation asimov”. Watch what happens to the relevance score of the top hit, and the scores of the other hits.

```
In [65]: req = requests.get("http://localhost:8983/solr/gettingstarted/select?q=foundation+asimov&wt=json")
        resp = req.json()
        for doc in resp['response']['docs']:
            print "%s - %s: %s" % (doc['score'], doc['id'], doc.get('name', '(no name)'))
```

```
2.0143478 - 0553293354: [u'Foundation']
0.009794351 - UTF8TEST: [u'Test with some UTF-8 encoded characters']
0.008570056 - SOLR1000: [u'Solr, the Enterprise Search Server']
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/spelling/suggest/tst
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-dataimporthandler-extras/deprecated-list
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-analytics/deprecated-list.html: (no name)
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clust
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/logging/log4j/packag
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-solrj/org/apache/solr/client/solrj/util/j
0.004328532 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-test-framework/org/apache/solr/analysis/
```

```
In [66]: explain = resp['debug']['explain']
        print explain['0553293354']
```

```
2.0143478 = (MATCH) sum of:
  0.05540521 = (MATCH) weight(_text:foundation in 1586) [DefaultSimilarity], result of:
    0.05540521 = score(doc=1586,freq=2.0), product of:
      0.14002547 = queryWeight, product of:
        1.1191508 = idf(docFreq=1785, maxDocs=2012)
        0.12511761 = queryNorm
      0.39567953 = fieldWeight in 1586, product of:
        1.4142135 = tf(freq=2.0), with freq of:
          2.0 = termFreq=2.0
        1.1191508 = idf(docFreq=1785, maxDocs=2012)
        0.25 = fieldNorm(doc=1586)
  1.9589427 = (MATCH) weight(_text:asimov in 1586) [DefaultSimilarity], result of:
    1.9589427 = score(doc=1586,freq=1.0), product of:
      0.99014795 = queryWeight, product of:
        7.9137373 = idf(docFreq=1, maxDocs=2012)
        0.12511761 = queryNorm
      1.9784343 = fieldWeight in 1586, product of:
        1.0 = tf(freq=1.0), with freq of:
          1.0 = termFreq=1.0
        7.9137373 = idf(docFreq=1, maxDocs=2012)
        0.25 = fieldNorm(doc=1586)
```

```
In [67]: print explain['SOLR1000']
```

```
0.008570056 = (MATCH) product of:
  0.017140113 = (MATCH) sum of:
    0.017140113 = (MATCH) weight(_text:foundation in 1577) [DefaultSimilarity], result of:
      0.017140113 = score(doc=1577,freq=1.0), product of:
        0.14002547 = queryWeight, product of:
          1.1191508 = idf(docFreq=1785, maxDocs=2012)
          0.12511761 = queryNorm
        0.122407116 = fieldWeight in 1577, product of:
```



```

    1.0 = tf(freq=1.0), with freq of:
        1.0 = termFreq=1.0
    1.1191508 = idf(docFreq=1785, maxDocs=2012)
    0.109375 = fieldNorm(doc=1577)
0.5 = coord(1/2)

```

Let's refine it a notch by using a specific field query for the author name.

```

In [68]: req = requests.get("http://localhost:8983/solr/gettingstarted/select?q=foundation+author:asimov")
        resp = req.json()
        for doc in resp['response']['docs']:
            print "%s - %s: %s" % (doc['score'], doc['id'], doc.get('name', '(no name)'))

0.02551029 - 0553293354: [u'Foundation']
0.00901925 - UTF8TEST: [u'Test with some UTF-8 encoded characters']
0.007891844 - SOLR1000: [u'Solr, the Enterprise Search Server']
0.004149459 - /Volumes/PGCD0803/etext03/vbgle11h/index.html: (no name)
0.003985983 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/spelling/suggest/tst
0.003985983 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-dataimporthandler-extras/deprecated-list
0.003985983 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-analytics/deprecated-list.html: (no name)
0.003985983 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-clustering/org/apache/solr/handler/clust
0.003985983 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-core/org/apache/solr/logging/log4j/packag
0.003985983 - /Users/dchud/Downloads/apps/solr-5.0.0/docs/solr-solrj/org/apache/solr/client/solrj/util/

```

What happened to the scores?

```

In [69]: explain = resp['debug']['explain']
        print explain['0553293354']

0.02551029 = (MATCH) product of:
    0.05102058 = (MATCH) sum of:
        0.05102058 = (MATCH) weight(_text:foundation in 1586) [DefaultSimilarity], result of:
            0.05102058 = score(doc=1586,freq=2.0), product of:
                0.1289442 = queryWeight, product of:
                    1.1191508 = idf(docFreq=1785, maxDocs=2012)
                    0.11521612 = queryNorm
                0.39567953 = fieldWeight in 1586, product of:
                    1.4142135 = tf(freq=2.0), with freq of:
                        2.0 = termFreq=2.0
                    1.1191508 = idf(docFreq=1785, maxDocs=2012)
                    0.25 = fieldNorm(doc=1586)
    0.5 = coord(1/2)

```

```

In [70]: print explain['SOLR1000']

0.007891844 = (MATCH) product of:
    0.015783688 = (MATCH) sum of:
        0.015783688 = (MATCH) weight(_text:foundation in 1577) [DefaultSimilarity], result of:
            0.015783688 = score(doc=1577,freq=1.0), product of:
                0.1289442 = queryWeight, product of:
                    1.1191508 = idf(docFreq=1785, maxDocs=2012)
                    0.11521612 = queryNorm
                0.122407116 = fieldWeight in 1577, product of:
                    1.0 = tf(freq=1.0), with freq of:
                        1.0 = termFreq=1.0
                    1.1191508 = idf(docFreq=1785, maxDocs=2012)
                    0.109375 = fieldNorm(doc=1577)
    0.5 = coord(1/2)

```

3.3 Faceting

This is only waving at the surface of what Solr can do, not even scratching it. Just as an example, here's what an result set faceted by author names looks like.

```
In [71]: req = requests.get("http://localhost:8983/solr/gettingstarted/select?q=*&facet=true&facet.fi
resp = req.json()
resp['responseHeader']
```

```
Out[71]: {'QTime': 89,
          'params': {'debugQuery': u'true',
                     'facet': u'true',
                     'facet.field': u'author',
                     'fl': u'*,score',
                     'indent': u'true',
                     'q': u'*:*',
                     'wt': u'json'}},
          'status': 0}
```

```
In [72]: resp['facet_counts']
```

```
Out[72]: {'facet_dates': {},
          'facet_fields': {'author': [u'George R.R. Martin',
3,
u'Lloyd Alexander',
2,
u'Rick Riordan',
2,
u'A NOVEL IN THREE PARTS BY FYODOR DOSTOEVSKY',
1,
u'Forster, E. M.',
1,
u'Glen Cook',
1,
u'Henry David Thoreau',
1,
u'Isaac Asimov',
1,
u'Jostein Gaarder',
1,
u'Michael McCandless',
1,
u'Orson Scott Card',
1,
u'Roger Zelazny',
1,
u'Rudyard Kipling',
1,
u'Sir James George Frazer',
1,
u'Steven Brust',
1,
u'gilevi',
1]}},
          'facet_intervals': {}},
```

```
u'facet_queries': {},  
u'facet_ranges': {}}
```

```
In [73]: resp['debug']['timing']
```

```
Out[73]: {u'prepare': {u'debug': {u'time': 0.0},  
    u'expand': {u'time': 0.0},  
    u'facet': {u'time': 0.0},  
    u'highlight': {u'time': 0.0},  
    u'mlt': {u'time': 0.0},  
    u'query': {u'time': 1.0},  
    u'stats': {u'time': 0.0},  
    u'time': 1.0},  
  u'process': {u'debug': {u'time': 1.0},  
    u'expand': {u'time': 0.0},  
    u'facet': {u'time': 86.0},  
    u'highlight': {u'time': 0.0},  
    u'mlt': {u'time': 0.0},  
    u'query': {u'time': 8.0},  
    u'stats': {u'time': 0.0},  
    u'time': 95.0},  
  u'time': 96.0}
```

86 milliseconds is pretty fast. Why is this so fast?